

BRAIN-COMPUTER INTERFACE (BCI)

Motor Imagery EEG Classification

Full End-to-End Engineering Pipeline: Preprocessing, EDA, Machine Learning & Web Deployment

Author / Team Lead: Data & Preprocessing Engineering • Cellula Team Project • August 2026

Project Overview & Mission Scope

Core Mission & Scope

- **Decoding Brain Motor Intention:** Classify whether a human is imagining LEFT or RIGHT hand movement strictly using scalp EEG signals (zero physical muscle actuation).
- **Raw Multi-Channel Dataset:** Processed 2,160 experimental trial CSV files captured via 4 scalp electrodes (FZ, C3, CZ, C4) at 250 Hz.
- **Multi-Stage Pipeline Engineering:** Designed, validated, and deployed the complete 4-stage pipeline: Preprocessing, Exploratory Data Analysis, Modeling, and Web App Deployment.
- **Assistive Technology Goal:** Targeted for neuro-rehabilitation and motorized assistive wheelchair/prosthetic control demos.

Delivered Engineering Assets

- ✓ **Clean Standardized Tensors:** X_clean.npy (0.5–40 Hz broadband) and X_clean_mubeta.npy (8–30 Hz SMR band) across 1,031 clean trials.
- ✓ **Comprehensive Audit Trail:** 100% transparent audit log (trial_audit_metadata.csv) tracking amplitude, variance, and rejection reasons for all 2,160 trials.
- ✓ **5 Evaluated Model Benchmarks:** Trained and cross-validated CSP+SVM, EEGNet, CNN, CNN-LSTM, and Transformer models.
- ✓ **Production Web App:** Interactive Flask web application (Deployment/app.py) with real-time CSV trial upload and safety rejection gates.

Neurophysiology of Motor Imagery (SMR Modulation)

1. Sensorimotor Rhythms

- Originates over the primary motor cortex (M1) and somatosensory strip.
- Mu Rhythm (8–12 Hz): Reflects idling cortical motor networks.
- Beta Rhythm (13–30 Hz): Associated with sensorimotor inhibition and active task execution.
- Electrodes C3 and C4 directly overlay the hand representation areas.

2. Contralateral ERD / ERS

- Event-Related Desynchronization (ERD): Imagining hand movement produces a localized power DECREASE in Mu/Beta bands.
- Contralateral Organization:
 - RIGHT Hand Imagery -> ERD on LEFT hemisphere (C3 electrode).
 - LEFT Hand Imagery -> ERD on RIGHT hemisphere (C4 electrode).
- Event-Related Synchronization (ERS): Power rebound post-imagery.

3. The Physical Challenge

- Extreme Low SNR: Cortical ERD is only 10–50 μV , heavily attenuated by skull bone and scalp tissue.
- Volume Conduction: Current spreads across the scalp, causing cross-electrode signal mixing.
- High Interference: Blinks (100–500 μV), muscle tension, and 50 Hz powerline hum swamp brain signals.
- Subject Variability: Individual alpha frequencies vary widely (8.5–12.5 Hz).

Raw Data Landscape & Hardware Reality

TOTAL RECORDED TRIALS

2,160

CSV files (1 trial/file)

ELECTRODE MONTAGE

4 Channels

FZ, C3, CZ, C4

SAMPLING FREQUENCY

250 Hz

2,500 samples (10.0s)

HARDWARE AMPLIFIER

ADS1299

24-bit ADC (Gain=24)

Hardware Calibration & Scale Factor

- **Raw 24-bit ADC Counts:** Raw CSV values were ~300,000 integer counts, reflecting raw Texas Instruments ADS1299 ADC output.
- **Microvolt Scale Constant:** $V = \text{Counts} \times (4.5V / (24 \times (2^{23} - 1))) \times 10^6 = \text{Counts} \times 0.02235174 \mu V/\text{count}$.
- **Restoring Physiology:** Multiplying raw counts by 0.02235 μV returned true physiological amplitudes (~20–100 μV).

Raw Data Flaws & Deficiencies

- ⚠ **Massive DC Offset & Drift:** Signals had huge half-cell DC offsets (~7,200 μV) that would cause filter edge ringing if unaddressed.
- ⚠ **50 Hz Powerline Intrusion:** Mains hum completely dominated high-frequency spectrum across all electrodes.
- ⚠ **Timestamp Jitter:** Inter-sample interval Δt ranged from 0.000005s to 0.040s; required uniform spline resampling to 250 Hz.

The 8-Stage EEG Preprocessing Pipeline

1. Ingestion & Scale

Convert 24-bit ADC counts to μV (0.02235). Spline resample to uniform 250 Hz.

2. Detrend & 50Hz Notch

Linear detrending removes DC drift. Zero-phase IIR Notch ($Q=30$) cuts powerline hum.

3. Bandpass Filtering

4th-order zero-phase Butterworth SOS (0.5–40 Hz broadband & 8–30 Hz Mu/Beta).

4. Spatial Re-reference

Common Average Reference (CAR: $V_i - \text{mean}$) cancels global environmental noise.

5. Artifact Rejection

Discard trials with PTP $> 200 \mu\text{V}$, dead channels ($\text{Var} < 0.5$), or variance $|Z| > 4.0$.

6. Baseline Subtraction

Subtract mean of pre-stimulus baseline window (0.0–0.5s) to center signal at 0 μV .


7. Task Epoch Window

Slice active motor imagery execution window (0.5–3.5s post-cue, exactly 751 points).

8. Conditioning & Scale

Soft-clip spikes via 4.5σ Winsorization. Standardize channels via Z-score scaling.


Quality Assurance & Artifact Rejection Audit



Artifact Audit & Cohort Metrics


- **Total Ingested Trials:** 2,160 raw files (100% evaluated)
- **Clean Retained Trials:** 1,031 trials (47.7% acceptance rate)
- **Corrupted Rejected Trials:** 1,129 trials (52.3% rejection rate)
- **Class Balance Maintained:** 533 Left (51.7%) vs 498 Right (48.3%) — Perfectly balanced 1.07 : 1 ratio
- **Extreme Amplitude Rejections:** 982 trials exceeded 200 μV PTP threshold (eye blinks, jaw clenching, motion)
- **Statistical Outlier Rejections:** 814 trials had robust variance $|Z| > 4.0$
- **Scientific Integrity:** Every rejected trial and its explicit reason is permanently logged in trial_audit_metadata.csv

EDA & Verification: Filtering & Spatial Decoupling



Spectral Verification Insights

- >40 dB attenuation of 50 Hz powerline interference without notch leakage.
- Clean preservation of Mu (8–12 Hz) and Beta (13–30 Hz) rhythm passband.
- Low-frequency drift rolloff eliminates half-cell electrochemical offsets.



CAR Spatial Decoupling Insights

- Before CAR: Channels were 93%–98% correlated due to common reference noise.
- After CAR: Cross-channel correlation dropped dramatically; C3 vs C4 correlation dropped to -0.27, unmasking independent hemispheric motor dynamics.

Modeling Strategy & Architecture Comparison

CSP + SVM

Classical Baseline

Common Spatial Patterns learns spatial variance filters in Mu/Beta band. Classified by Linear SVM.

Features: Log-Variance
Parameters: Minimal
Speed: 1.35s

EEGNet

Specialized EEG DL

Compact CNN with temporal conv (1x64), depthwise spatial conv (4x1), separable conv, and ELU.

Features: Raw Epoched
Parameters: 1,426
Speed: 15.57s

CNN

2D Convolutional

Multi-stage temporal-spatial convolutions with feature flattening and dense classification head.

Features: Raw Epoched
Parameters: 26,834
Speed: 18.56s

CNN-LSTM

Spatio-Temporal

Cascades convolutional feature maps into a recurrent LSTM (hidden=32) for temporal sequence memory.

Features: Raw Epoched
Parameters: 37,202
Speed: 23.92s

Transformer

Self-Attention

Temporal patch embedding, learnable 1D positional encoding, and 2-layer 4-head Transformer Encoder.

Features: Raw Epoched
Parameters: 44,402
Speed: 15.45s

Experimental Results: The Performance Reality

CRITICAL RESULT: All evaluated models achieved ~50% to 52% test accuracy (chance-level binary performance). Permutation tests confirm that this is indistinguishable from random guessing.

Model	Input Band	Parameters	Accuracy	Balanced Acc	F1 Score	ROC-AUC	Train Time
CSP + SVM	8–30 Hz	Minimal	50.15% ± 0.7%	49.65%	0.3866	0.5136	1.35s
EEGNet (Best)	0.5–40 Hz	1,426	51.79% ± 1.6%	51.34%	0.3447	0.5170	15.57s
CNN	0.5–40 Hz	26,834	50.72% ± 2.5%	50.40%	0.4221	0.4917	18.56s
CNN-LSTM	0.5–40 Hz	37,202	50.92% ± 1.3%	49.88%	0.2033	0.4887	23.92s
Transformer	0.5–40 Hz	44,402	50.34% ± 2.4%	49.40%	0.2831	0.4908	15.45s
Shuffled Baseline	Permuted	—	50.50% ± 1.1%	50.00%	—	0.5000	—

Key Diagnostic Finding

- Zero Pipeline Leakage: Strict 5-Fold Stratified CV with isolated fold transformations ensures metrics reflect genuine out-of-fold generalization.
- Permutation Test Proof: 20 iterations of shuffled labels yielded 50.50% accuracy; the real-label CSP score of 50.15% is statistically identical to random guessing ($p = 0.55$).
- Provenance Verification: Label alignment audits confirmed 100% data and label synchronization with zero indexing errors.

Why Are Results Not Good? Part I: Spatial & Algorithmic

1. Extreme Spatial Sparsity

- Benchmark BCI datasets (e.g. BCI Comp IV 2a) use 22 to 64 electrodes.
- With only 4 channels (FZ, C3, CZ, C4), the spatial covariance matrix is only 4x4, offering just 4 degrees of freedom.
- Common Spatial Patterns (CSP) mathematically requires redundant spatial dimensions to compute spatial filters that cancel volume conduction.
- Without surrounding Laplacian rings around C3/C4 (e.g. FC3, CP3, C1, C5), local sensorimotor activity cannot be isolated.

2. CAR Distorts Contrast

- Common Average Reference subtracts the mean across all electrodes: $V_C3 - (FZ+C3+CZ+C4)/4$.
- Expanding this yields: $(3/4)C3 - (1/4)C4 - (1/4)FZ - (1/4)CZ$.
- Because C3 and C4 are the contralateral motor channels, CAR forces 25% of C4's signal directly into C3!
- With only 4 channels, CAR inadvertently washes out the lateralized difference between Left and Right hand mental imagery.

3. Variance Signal Flattening

- Normalizing each channel of every trial to unit variance (Z-score) forces channel variance to exactly 1.000.
- Diagnostic audit in `class_signal_summary.csv` verified that variance and RMS are identically 1.0 across all trials.
- CSP and ERD are fundamentally VARIANCE-BASED mechanisms (detecting power drops in C3 relative to C4).
- Forcing every channel to unit variance flattens the amplitude discrepancy classifiers rely on.

Why Are Results Not Good? Part II: Biological & Experimental

4. Cross-Subject Pooling

- Trials were pooled across multiple subjects without subject identifier metadata.
- Individual Alpha Frequency (IAF) varies widely between 8.5 Hz and 12.5 Hz across humans.
- Anatomical hand knob locations differ across skulls.
- Training a single global model across pooled subjects without domain adaptation washes out individual ERDs into the group noise floor.

5. BCI Illiteracy (15–30%)

- In clinical BCI literature, 15% to 30% of human subjects are 'BCI illiterate'.
- Untrained participants fail to produce measurable Mu/Beta desynchronization during imagery without neurofeedback training.
- Without live sensory feedback to condition the brain, naive subjects' mental imagery produces zero detectable voltage change.

6. Severe Acquisition Noise

- 52.3% of trials (1,129 out of 2,160) had to be rejected due to massive voltage bursts ($>200 \mu\text{V}$) or dead flatlines.
- This extraordinarily high contamination rate indicates poor skin-electrode contact impedance and excessive subject motion.
- Even retained clean trials likely contained sub-threshold muscle micro-tension that obscured faint cognitive brain rhythms.

Web Deployment Architecture & Clinical Safety

Production Flask REST Application

- **Lightweight Web API:** Engineered with Flask (Deployment/app.py) for full control over async JavaScript fetch endpoints and UI state management.
- **Dual Ingestion Modes:**
 - Mode 1: Drag-and-drop CSV trial upload.
 - Mode 2: One-click random clean example evaluation.
- **End-to-End Inference:** Invokes `transform_single_trial()` on the raw CSV, applying identical calibration, detrending, notch, bandpass, CAR, and epoching in real time.
- **Ultralight Model Serving:** Serves PyTorch EEGNet (1,426 parameters, <20 KB weight file), executing inference in under 5 milliseconds on CPU.

Clinical Safety Gate Mechanism

- ✓ **The Biomedical Risk:** In assistive robotics (e.g. wheelchair control), predicting on noisy EEG contaminated by a cough or blink could cause catastrophic false actuations.
- ✓ **Production Safety Gate:** Before executing neural network forward pass, the trial is audited against rejection thresholds ($PTP \leq 200 \mu V$, $Var \geq 0.5$).
- ✓ **Rejection Fail-Safe:** If corrupted, the app REFUSES to output a prediction. It halts immediately and returns an explicit warning alert:
- ✓ **Response Payload:** `{"is_valid": false, "rejection_reasons": "extreme_ptp_amplitude_C3_(485.2uV > 200.0uV)"}`

Roadmap to High Performance (>80% Accuracy)

1. Pre-Processing

- Replace per-trial Z-score with pre-cue baseline normalization (preserves spatial power contrast).
- Replace 4-channel CAR with localized bipolar derivations (C3 - CZ, C4 - CZ).
- Apply Independent Component Analysis (ICA) to remove blinks/EMG instead of discarding 52% of trials.
- Tune filter passbands to Individual Alpha Frequency (IAF).

2. Machine Learning

- Implement Filter Bank CSP (FBCSP) across 4–40 Hz sub-bands with feature selection.
- Apply Riemannian Geometry on covariance matrices (Tangent Space Logistic Regression).
- Use Domain Adaptation (CORAL / DANN) to align cross-subject distributions.
- Expand dataset with EEG time-shift and Mixup augmentation.

3. Hardware & Montage

- Increase electrode density from 4 to 16–32 channels focused on sensorimotor strip.
- Deploy true surface Laplacian rings around C3 and C4 (FC3, CP3, C1, C5).
- Use active shielded wet electrodes; enforce skin impedance < 5 k Ω .
- Record explicit Subject ID and Session ID metadata.

4. Experimental Protocol

- Implement real-time visual neurofeedback to condition subject motor imagery.
- Train subjects over 3–5 sessions to overcome BCI illiteracy.
- Synchronize precise cue onset timestamps using hardware triggers.
- Deploy confidence-gated thresholds in production web demos.

Conclusions & Engineering Takeaways

Engineering Achievements

- ✓ **Modular Production Pipeline:** Successfully built, tested, and validated an 8-stage zero-phase preprocessing pipeline from raw ADC counts to normalized tensors.
- ✓ **Scientific Integrity & Audit:** Provided 100% transparent audit metadata for every trial, verifying label alignment and provenance.
- ✓ **Extensive Benchmark Suite:** Evaluated 5 competitive architectures with strict Stratified 5-Fold Cross-Validation, eliminating data leakage.
- ✓ **Full Deployment Ready:** Delivered a lightweight Flask web application with live CSV prediction and clinical artifact rejection safety.

Scientific & Diagnostic Takeaways

- 💡 **Honest Empirical Science:** The ~50% baseline accuracy is not an engineering failure, but a proven consequence of 4-channel spatial limits, CAR contrast mixing, and pooled inter-subject variability.
- 💡 **Garbage In, Garbage Out:** Advanced deep learning (Transformers, CNN-LSTM) cannot compensate for raw data where physiological class differences are absent ($p > 0.58$).
- 💡 **Clear Path to High Accuracy:** The clear technical roadmap (FBCSP, Riemannian geometry, local bipolar derivations, and dense electrode arrays) provides a structured blueprint for future 80%+ BCI systems.